

1. Analysis

How is the memory subsystem structured on your machine. On Linux system likwid-topology (developed by Google) is a good tool to know how the memory system is architected.

Question: How much DRAM memory? Memory speed? Memory Technology? (Not all is exposed unless you are root.)

-> DRAM memory : 65314.3 (Domain 0) + 65536 MB (Domain 1)

-> Memory Speed : DDR4 speed 1600/1866/2133

(https://ark.intel.com/products/83361/Intel-Xeon-Processor-E5-2667-v3-20M-Cache-3_20-GHz)

-> Memory Technology : DDR4

Question: How much bandwidth can your processor draw from the bus? You will need to discover width of the memory bus and clock speed of the memory bus. Full duplex/half duplex? (On intel processor, ARK often indicates something.)

-> Bandwidth processor can draw from bus : 9.6 GT/s (GigaTransfer per second) QPI

(https://ark.intel.com/products/83361/Intel-Xeon-Processor-E5-2667-v3-20M-Cache-3_20-GHz)

-> Memory Bandwidth mentioned on arch.intel.com site : 68 GB/s.

-> Xeon Processor E5 2667 v3 is full duplex.

Question: What is the highest level of data cache? How big is it? How is it shared across core? Same question for all levels of data cache.

-> Highest level of cache is L3 cache. It's size is 20MB and even processors share one L3 cache and odd number processors share another L3 cache.

-> L2 cache is of size 256KB and following groups share the cache.

- (0 2)

- (4 6)

- (8 10)

- (12 14)

- (1 3)

- (5 7)

- (9 11)

- (13 15)

-> L1 cache is of size 32KB and is share in same group as L2 cache is shared.

2. Bandwidth

Question: Write a code to measure read bandwidth

-> Attached

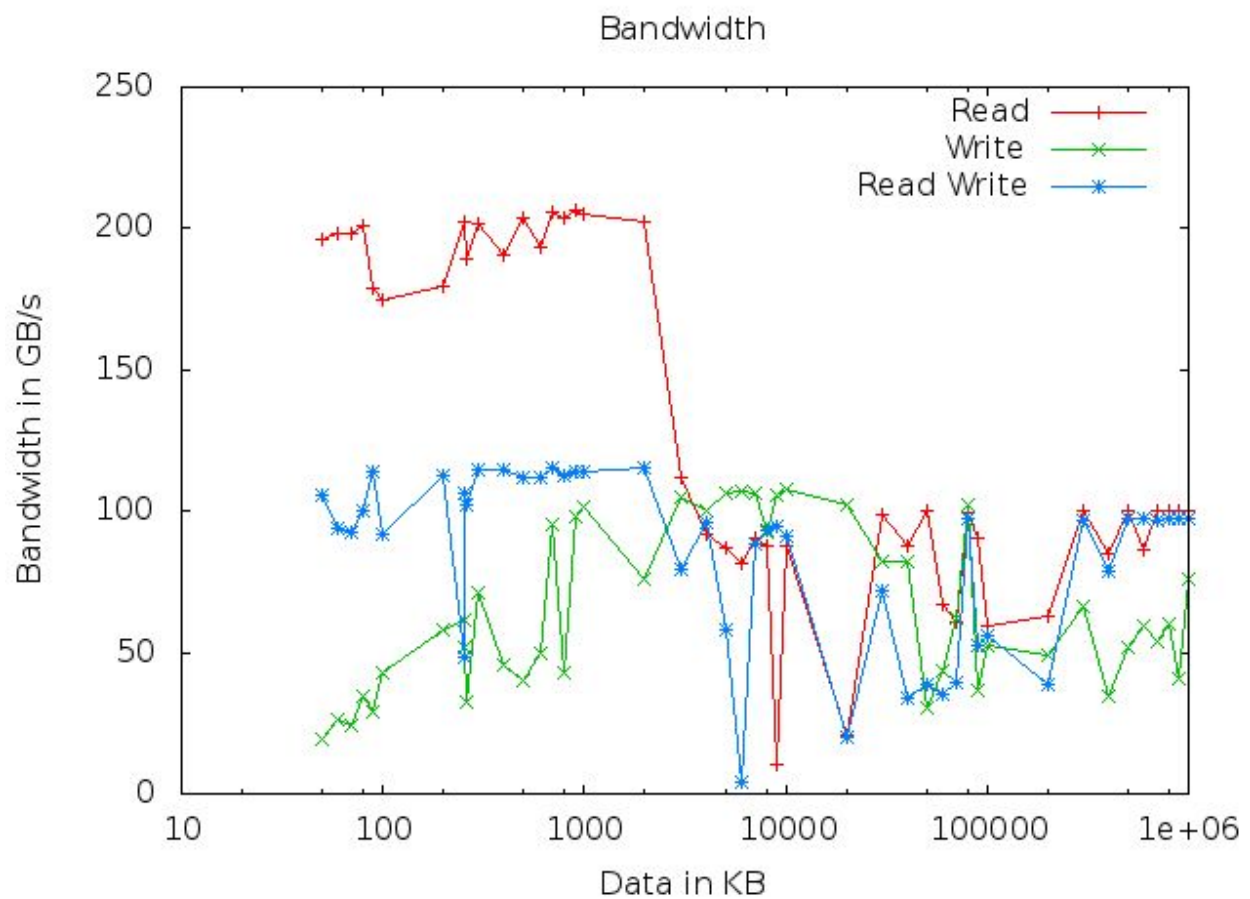
Question: Write a code to measure write bandwidth

-> Attached

Question: Write a code to measure read/write bandwidth

-> Attached

Question: Measure read, write, read/write bandwidth at different size of data. (From 1KB to 200MB)



Data points in Bandwidth.txt file

3. Latency

Question: Write a code that traverses some element of a linked list as described above.

-> Attached

Question: Write function to create a linked list which structure is agreeable to the core structure.

-> Wrote code such that, the i th element of list holds the address of $i+1$ element

Question: What is the latency in that case?

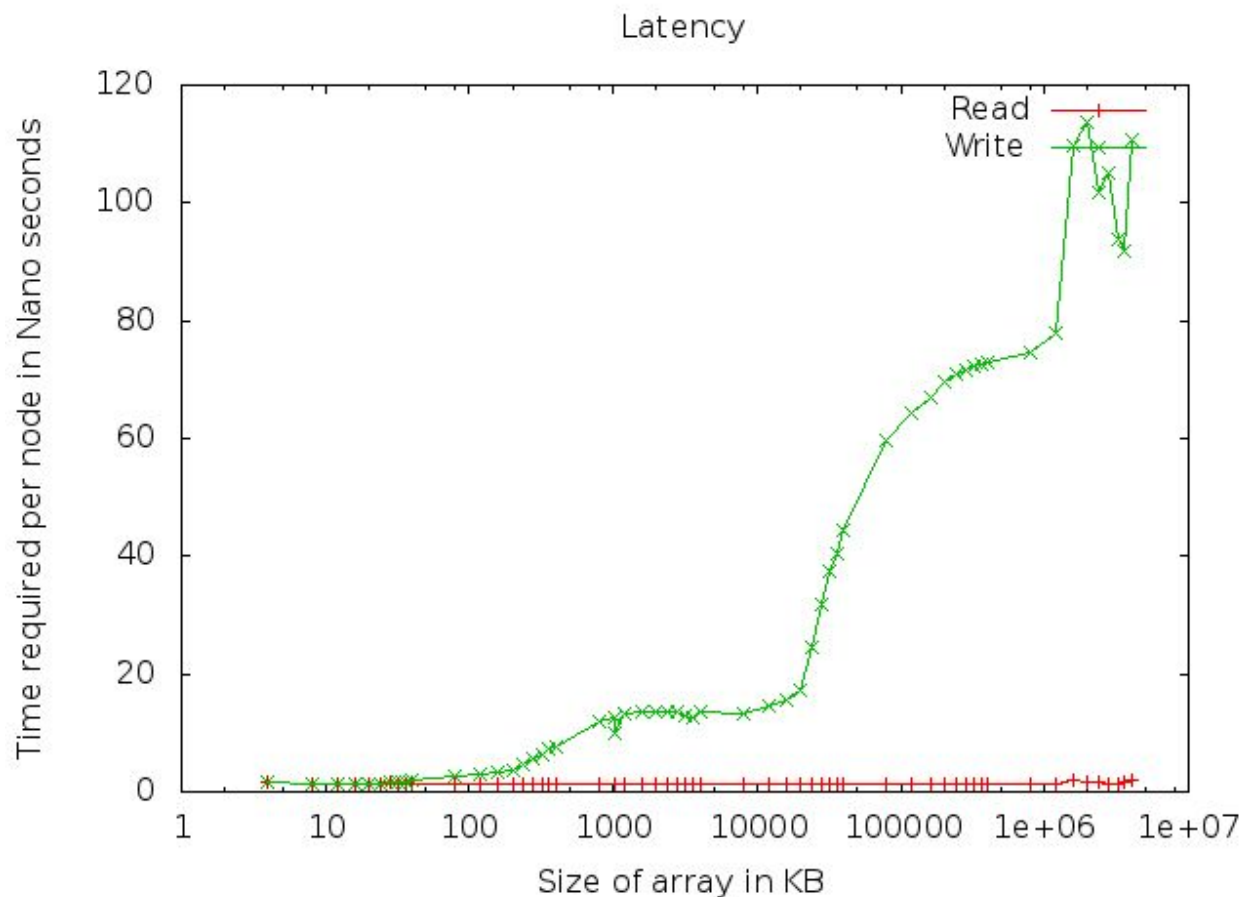
-> 1 nano second, as hardware prefetching occurs

Question: Write a function to create a linked list that will jump around the memory subsystem while avoiding to build short cycles.

-> Scrambled the array index, using another array.

Question: For different size (From 1kB to 200MB), what is the latency of the system?

->



Red line is Sequential

Green Line is Scattered

Data points in latency.txt